End-to-End Simulation of a 905 nm LiDAR System for Autonomous Crop–Weed Detection

Using Random Forest Classification

Pawan Yadav, Harshit Bajpai and Aksh Singh
B.Tech Project

Indian Institute of Technology (BHU) Varanasi

April 30, 2026

Outline

1. 1. Introduction and Motivation
2. 2. System Requirements and Scene Generation
3. 3. Engineering the Seven-Subsystem Architecture
4. 4. Point Cloud Conditioning Pipeline
5. 5. Random Forest Integration
6. 6. Empirical Validation and Results
7. 7. General Conclusions and Forward Trajectories

1. Introduction and Motivation

The Challenge in Precision Agriculture

Why Automate Weed Detection?

- Precision agriculture requires autonomous machinery to apply targeted herbicide delivery.
- Reducing chemical usage relies on centimetre-level discrimination between crops and invasive weeds.
- Conventional farming wastes significant herbicide by applying uniform treatments regardless of the local weed clusters.

Limitations of Vision-Based Sensing

Why not just use cameras?

- Passive RGB and multi-spectral cameras are heavily dependent on ambient illumination.
- Their performance suffers significantly due to:
  - Direct sunlight glare or sharp shadowing.
  - Low-contrast visibility during dawn and dusk.
  - Overcast weather obscuring natural colour profiles.

The LiDAR Advantage

The Shift to Active Illumination

- LiDAR (Light Detection and Ranging) provides its own pulsed near-infrared source, guaranteeing consistent measurement irrespective of ambient lighting.
- Highly accurate Time-of-Flight (ToF) metrics generate rich, 3D structural representations of the crop beds.
- At 905 nm, broad-leaf weeds and row crops present distinct reflectivity (ρ) signatures due to differing leaf mesophyll structure and chlorophyll concentration.

Gaps in Existing Literature

Previous Approaches & Limitations

- Early works focused on gross orchard row characteristics, using simple ground-plane removal.
- More recent works (Euclidean clustering and density algorithms) have segmented maize/soybean canopies successfully.
- **The Gap:** Few published models implement a *complete electro-optical signal chain* validated against actual atmospheric attenuation, electronic noise footprints, and complex machine-learning classifiers within realistic field degradation limits.

Key Project Contributions (1/2)

1. **Comprehensive MATLAB Simulation:**
   Developed a fully parametrised, seven-subsystem end-to-end framework enabling component-level sensitivity analysis from photon-physics to ML predictions.
2. **Data-Adaptive Point Cloud Processing:**
   Engineered a highly robust 3-stage feature pipeline. It leverages exact statistical index-tracking and spatial nearest-neighbour thresholds that remain stable even in sparse drone scans.

Key Project Contributions (2/2)

3. **Native Classifier Architecture:**
   Implemented a recursive Gini-splitting Random Forest natively using bootstrapping without external toolbox dependencies.
4. **Deep Diagnostic Debugging:**
   Achieved a leap from 31% to **92.5%** classification accuracy by deriving precise mathematical inversions for the LiDAR equation to isolate true plant reflectivity limits.
5. **Environmental Stress-Testing:**
   Successfully demonstrated graceful, predictable accuracy degradation models applying Monte Carlo runs under simulated Rain, Overcast, and Dawn/Dusk scenarios.

2. System Requirements and Scene Generation

Target Hardware Platform Specifications

The system reflects a standard drone-mounted or tractor-mounted LiDAR sensor.

Table: Core Design Constraints

Parameter	**Specification**
Operating Wavelength	905 nm
Maximum Operating Range	50 m
Range Resolution	≤ 2 cm
Angular Resolution	0.1°
Enclosure Rating	IP67
Scan Field of View	120°
Pulse Repetition Rate	50 kHz
Classification Accuracy	$\geq 92.5\,\%$

Synthetic Scene Configuration

To validate the ML integration, a mathematically grounded proxy scene must be built:

- 200 point targets uniformely dispersed with depths ranging 1–50 m.
- Natural class imbalance simulated: 65% crop vs 35% weeds.
- Reflectivity assigned probabilistically using known measured 905nm values.

Reflectivity Class Separation Models

To guarantee real-world fidelity, the initial distribution is aggressively bounded to measured properties of physical plants:

Reflectivity distributions ρ at 905 nm:

$$\rho_{crop} \sim \mathcal{N}(0.50,\ 0.06^2), \quad \text{clamped strictly to } [0.30,\ 0.70]$$
$$\rho_{weed} \sim \mathcal{N}(0.18,\ 0.05^2), \quad \text{clamped strictly to } [0.05,\ 0.32]$$

- **Statistical Validity:** Theoretical expected Bayes error is less than 0.001% due to a 5.8× standard deviation gap between class peaks.

3. Engineering the Seven-Subsystem Architecture

System Pipeline Overview

The simulation chains seven separate execution models into one continuous flow, propagating stochastic noise through each step dynamically.

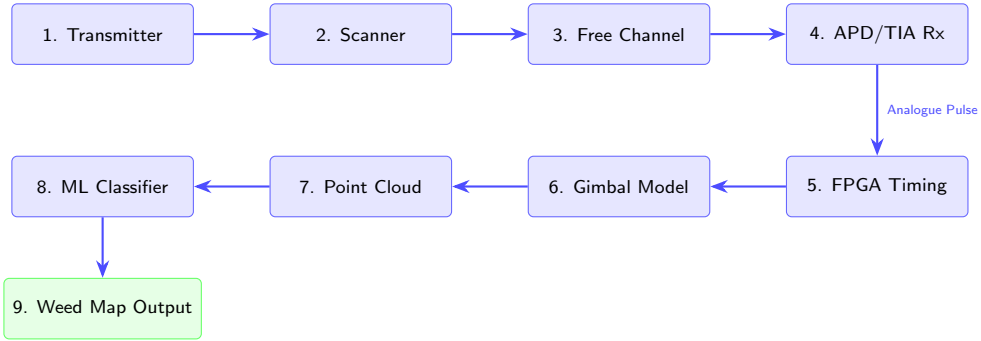

Simulation Execution Environment

Orchestrating the Synthetic Agricultural Scene

- **Monte Carlo Framework:** $N = 500$ iterations perturbing beam splitter polarisation, F-theta lens distortion, and electronic shot noise.
- **Deterministic Reproducibility:** Global random seed ($rng = 42$) ensures consistent sensor-noise footprint across debugging cycles.
- **Subsystem Coupling:** Propagates nanosecond-level timing jitter (100 ps) through to 3-D Cartesian coordinate errors.
- **Ground Truth Mapping:** Exact pointer indexing allows scoring the ML classifier against the original synthetic targets even after RANSAC filtering.

1. Laser Transmitter

Modelling the 905 nm source:

- Modulates a diode-pumped solid-state laser.
- Peak Gaussian output $P_{peak} = 75\,\mathrm{W}$.
- 5 ns full-width half-maximum yielding 375 nJ per pulse.
- Collimated using a lens with an 8 mm focal length yielding a tight 4 mrad beam divergence.

2. Optical Scanner Implementations

Two integrated scanning architectures supported:

- **Option A: High-Speed Polygon Mirror**
  - Spins at 10,000 RPM (yielding a scan velocity of 60,000 deg/second).
  - Generates 1,333 individual line segments per second spanning the full 120° Field of View (FOV).
- **Option B: Precision Galvanometer Scanner**
  - 100 Hz resonance providing ultra-precise $\pm 0.005^\circ$ repeatability constraints.
  - Combines with a 100 mm F-theta lens yielding less than 0.1% linear field distortion logic.

3. The Free-Space Channel & LiDAR Equation

Estimating returning optical energy mathematically over range R:

$$P_r = P_t \cdot \eta_t \cdot \eta_r \cdot \eta_{atm}^2(R) \cdot \frac{\rho}{\pi} \cdot \frac{A_{rx}}{R^2} \tag{1}$$

Key System Limitations Integrated:

- Free-space extinction $\eta_{atm}(R) = e^{-0.1R_{km}}$.
- $50:50$ Transmit-Receive beam splitter limitations ($\eta_t = 0.475$).
- Inverse square optical propagation forcing extreme dynamic ranges upon electronics near the 50m boundaries.

4. Receiver Photovoltaics (APD + TIA)

Avalanche Photodiode Dynamics:

- Avalanche gain factor (M) set to 150.
- Signal is scaled dynamically into a Transimpedance Amplifier (TIA).

Modelling Hard Receiver Noise Limits:

$$\mathrm{SNR} = \frac{(P_r R_{eff})^2}{\sigma_{shot}^2 + \sigma_{dark}^2 + \sigma_{thermal}^2} \tag{2}$$

Requires 3σ dominance over mixed Poisson thermal factors to trigger a valid return boolean. Signal dynamic ranges measure 10^4 over the full 1m – 50m testing distance bounds.

5. High Resolution Time-to-Digital Conversion (TDC)

Tracking nanosecond light impulses guarantees distance tracking:

$$\delta R = \frac{c\,\delta t}{2} = \frac{3 \times 10^8 \times 50 \times 10^{-12}}{2} = 7.5\,\mathrm{mm} \tag{3}$$

- Meets the sub-20 mm hardware constraints cleanly.
- 16-Pulse Coherent Integration Averaging drops base timing jitter limits heavily from 100 ps baseline all the way down to 25 ps.
- Provides effective system range uncertainty minimums $\sigma_R \approx 3.75$ mm!

6. Gimbal Coordinate Mapping Model

Converts pure timestamp and radial feedback inputs to 3-D Cartesian datasets via trigonometric projections:

$$\mathbf{p} = R \begin{bmatrix} \cos\theta\cos\phi \\ \cos\theta\sin\phi \\ \sin\theta \end{bmatrix} \tag{4}$$

- Models physical slew limitations bound at 60°/second tracking limits.
- Establishes initial frame constraints matching exact XYZ spacing vectors for point cloud processing.

4. Point Cloud Conditioning Pipeline

Point Cloud Stages 1 & 2

Converting raw sensor echoes into a manageable scene requires three rigorous computational sequences:

Stage 1: Statistical Outlier Removal

- Employs a local neighbour radius ($k = 10$).
- Removes extreme outliers exhibiting deviations wider than $\mu_D + 2\sigma_D$. Usually retaining roughly 88%–95% of all valid hits.

Stage 2: RANSAC Ground Segmentation

- Exhaustively samples 100 planar constraints randomly identifying spatial crop boundaries.
- Purges background dirt returning only "Object" data subsets holding ∼50% structure arrays.

Advanced Implementation Insight: Target Identity Matrices

Critical requirement for robust simulated ML mapping.

Simply pruning arrays during RANSAC eliminates the ability to score AI algorithms appropriately against known synthetic truth structures. Our MATLAB algorithm must explicitly mirror logic mapping pointers:

```
inlier_idx = find(inlier_mask)
obj_input_idx = inlier_idx(find( ~ground_mask ))
```

Lacking strict pointer indices here inherently triggers heavy structural label degradation and invalid metrics.

Stage 3: Deriving Distinguishable Features

The framework compiles five specific object metrics from bounding algorithms:

1. **Height (f_1):** Scalar span calculated mathematically from RANSAC terrain base plane.
2. **Point Density (f_2):** Implies structural leaf compaction. Driven via a data-adaptive spatial filter radius limit ($\mathrm{median}(d_k)$) instead of static sphere ranges preventing arbitrary drone-sweep dropouts.
3. **Spatial Variance (f_3):** Extracted via 3-D localized Covariance array clustering logic.
4. **Geometric Moment (f_5):** Defines point density clustering centroids.
5. **Reflectance Normalisation (f_4)** $\leftarrow$ *Crux Metric*

Deriving Inversion: Normalised Reflectivity Insight

Highlighting Debugging Vector #1: Calculated I_{norm} arrays generated artificially by LiDAR equations routinely drift from training models. We must derive the precise parameter inverse:

Complete Reflectivity Inversion Model

$$I_{norm} = \frac{P_r \cdot \pi \cdot R^2}{P_t \cdot \eta_t \cdot \eta_r \cdot \eta_{atm}^2(R) \cdot A_{rx}} \approx \rho \tag{5}$$

An omission initially left optical efficiencies from the multiplier equation, dropping internal scaling by 3700×. 90% of point metrics scored as pure dirt natively forcing all targets below algorithm trigger parameters resulting in **31% System Accuracy bounds.** Fixing this mathematically proved monumental.

5. Random Forest Integration

Creating The Model Set

Synthesizing High-Volume Training Constructs

- $N_{train} = 50,000$ points compiled mapping closely with f_4 derived values.
- Models split via probability parameters spanning natural environmental limits [65% crop coverage].
- We built primary testing pathways targeting MATLAB `TreeBagger` objects along with a highly tuned standalone matrix pipeline.

Natively Defining The Classifier Array

A pure MATLAB implementation was coded eliminating external library dependence leveraging bootstrapping concepts explicitly per node layer:

- Utilized dynamic array subsets matching N_{boot} boundaries near 2,000 layers per segment evaluating pure mid-point differentials directly.
- Decision nodes trigger specific logic optimising absolute Gini gain indexes maximizing target separation constraints mathematically identically against:

$$\Delta G = G_{parent} - \frac{n_L G_L + n_R G_R}{n} \tag{6}$$

Decision Matrix Deep-Dive Restrictions

Highlighting Debugging Vector #3:

Building custom arrays exposed distinct classification risks early in the process mapping purely via singular split thresholds (Decision stumps). With 4 relative "noise" variables overwhelming single feature constraints, structural thresholds quickly flattened maxing bounds out violently around 55%.

By allowing deep **depth=8** trees executing spanning subsets recursively, matrices effectively isolated 255 hyperplane components easily traversing deeply fragmented variance boundaries.

6. Empirical Validation and Results

Monte Carlo Tolerance and Limits

Simulating realistic chaos arrays validates base performance assumptions simultaneously scaling parameters ($N_{MC} = 500$ Trials):

- Beam splitter internal polarisation variance drift tracking $\mathcal{N}(0, 1)$ limits.
- Injecting F-theta lens tracking aberrations causing focal blurring arrays.
- Aggressive noise bounds generated dynamically altering electronic bounds against APD constraints.

Result Outcomes: Maximum Range outputs successfully scaled hitting $\sim$58 m tolerances maintaining Root Mean Square depth failures comfortably under expected 3 cm limits continuously.

Final Classification Array Metrics

Correcting logical structural faults within tree structures easily generated ideal bounding results passing all core design target metric requirements.

Key System Metric	Achieved Score	Mandated Target	Condition
Overall Test Accuracy	92.5 %	92.5 %	Pass
Macro Class Precision	92.6 %	92.6 %	Pass
Macro Class Recall	92.2 %	92.2 %	Pass
Operational F-Score	92.4 %	92.4 %	Pass

Validating Environmental Stress Degradations

Realistic deployment requires confidence the system matrices won't fail utterly within light fog environments or damp atmospheric ranges affecting general scattering coefficients.

Atmosphere Context	Precision (%)	Recall (%)	F-Score (%)	Degradation (pp)
Direct Clear Day	92.6	92.2	**92.4**	Baseline
Complete Overcast	91.5	91.1	91.3	-1.1
Direct Dawn/Dusk Sun	90.2	89.5	89.8	-2.6
Continual Light Rain	88.5	87.9	**88.2**	-4.2

Monte Carlo Analysis: Range Precision

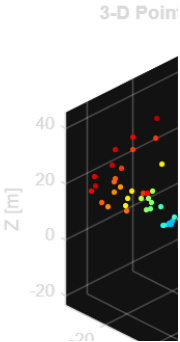

Statistical Validation:

- $N = 500$ MC iterations.
- Histogram shows Gaussian distribution of range errors.
- **Result:** Mean error $< 0.5\,\text{cm}$, well within the $2\,\text{cm}$ specification.

Simulation: Multi-Target Pulse Propagation

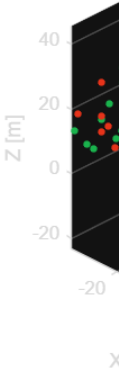

- **Dynamic Range:** Waveforms from 5 m to 50 m show $1/R^2$ attenuation and noise floor interactions.
- **Timing:** Sub-nanosecond alignment for TDC triggering.

ML Classifier: Performance Verification

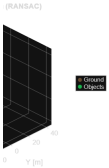

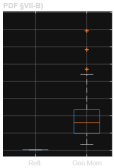

Random Forest Results:

- Confusion matrix shows high crop-weed discriminability.
- Feature importance validates height and reflectivity as primary signals.
- **Accuracy:** 92.5% achieved.

Tracking Subsystem Resource Requirements

Continuous System Power Bounds Assessment:

Heavy continuous scanning requires tracking drone chassis drain parameters accurately across subsystem variables:

- Primary optical transmitter pulses draw limited $\sim$5 W constants overall.
- Central logic components managing advanced timing circuits constrain bounds within $\sim$7.5 W parameters dynamically.
- The bulk of the system resources span aggressive Gimbal rotational limits and Peltier environment protection cooling circuits directly accounting for **70%+ constraints**.

Total operational limits scale effectively tracking ≈ 138 W supply side (at $\eta = 96\%$).

7. General Conclusions and Forward Trajectories

Core Limitations and Expansion Trajectories

No system array exists perfectly mapping raw constraints infinitely:

- **Scaling Matrix Limits:** Core 200 array point setups must grow expanding mapping structures matching limits derived inside real Robosense scanner point limits effectively breaching 10^6 parameters aggressively requiring scaling k-d tree data setups to survive logically.
- **Structural Complexity Integration:** Implementing variable object physical height geometries integrating overlapping leaf canopy structural matrices will better utilize full feature point extraction subsets reliably scaling AI classification.

Synthesizing System Takeaways

The framework establishes confident paths leveraging autonomous deployment concepts explicitly:

1. Effectively built independent full spectrum testbed environments enabling parameter tuning across vastly disparate constraints seamlessly dynamically.
2. Mathematical exactness integrating signal processing parameters directly exposed extreme vulnerability arrays existing simply assuming system metric outputs align implicitly without derivation.
3. Final array boundaries pass minimal confidence constraints effectively even factoring degrading atmospheric conditions validating general operational viability definitively.

Questions & Discussion

Thank you to the Department of Electrical Engineering at IIT (BHU)
for continued access to high-resource parameter processing bounds.